Supplementary Materials

Task-level influence-score distributions at nominal poisoning 1%

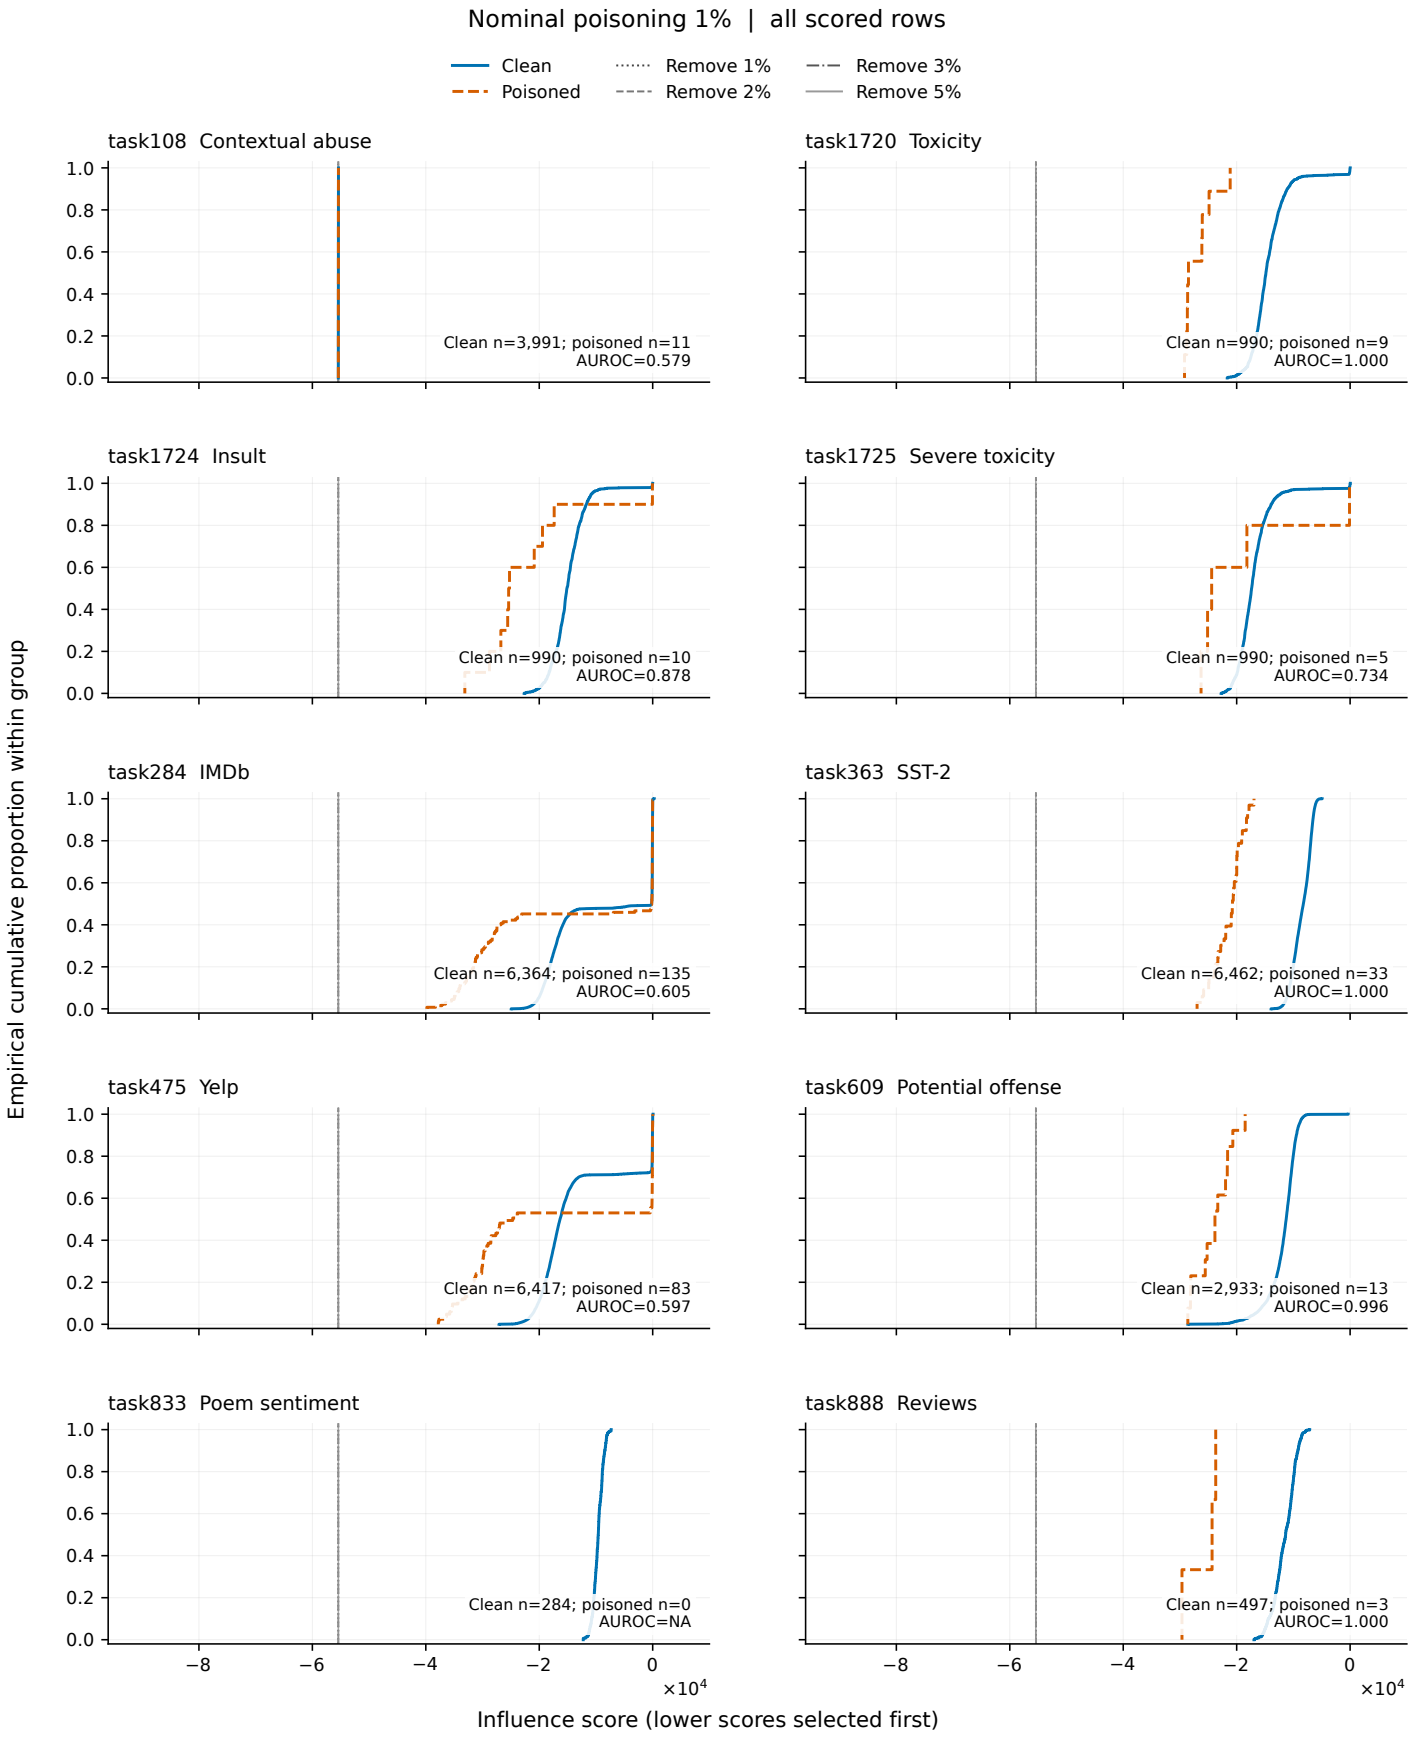


Figure S1. Nominal poisoning 1%. Empirical cumulative distributions of clean and poisoned scores within each source task. Curves are normalized within each task and group. Vertical lines mark global removal budgets of 1%, 2%, 3% and 5% of 30,220 scored rows. Lower scores are selected first. All scored rows, including references, are included. AUROC uses negative influence scores with half credit for ties. Task833 has no poisoned examples; task888 has very few. Near-coincident cutoffs may appear as one line. Scores from different adapters and reference sets are not calibrated across poisoning rates.

Supplementary Materials

Task-level influence-score distributions at nominal poisoning 3%

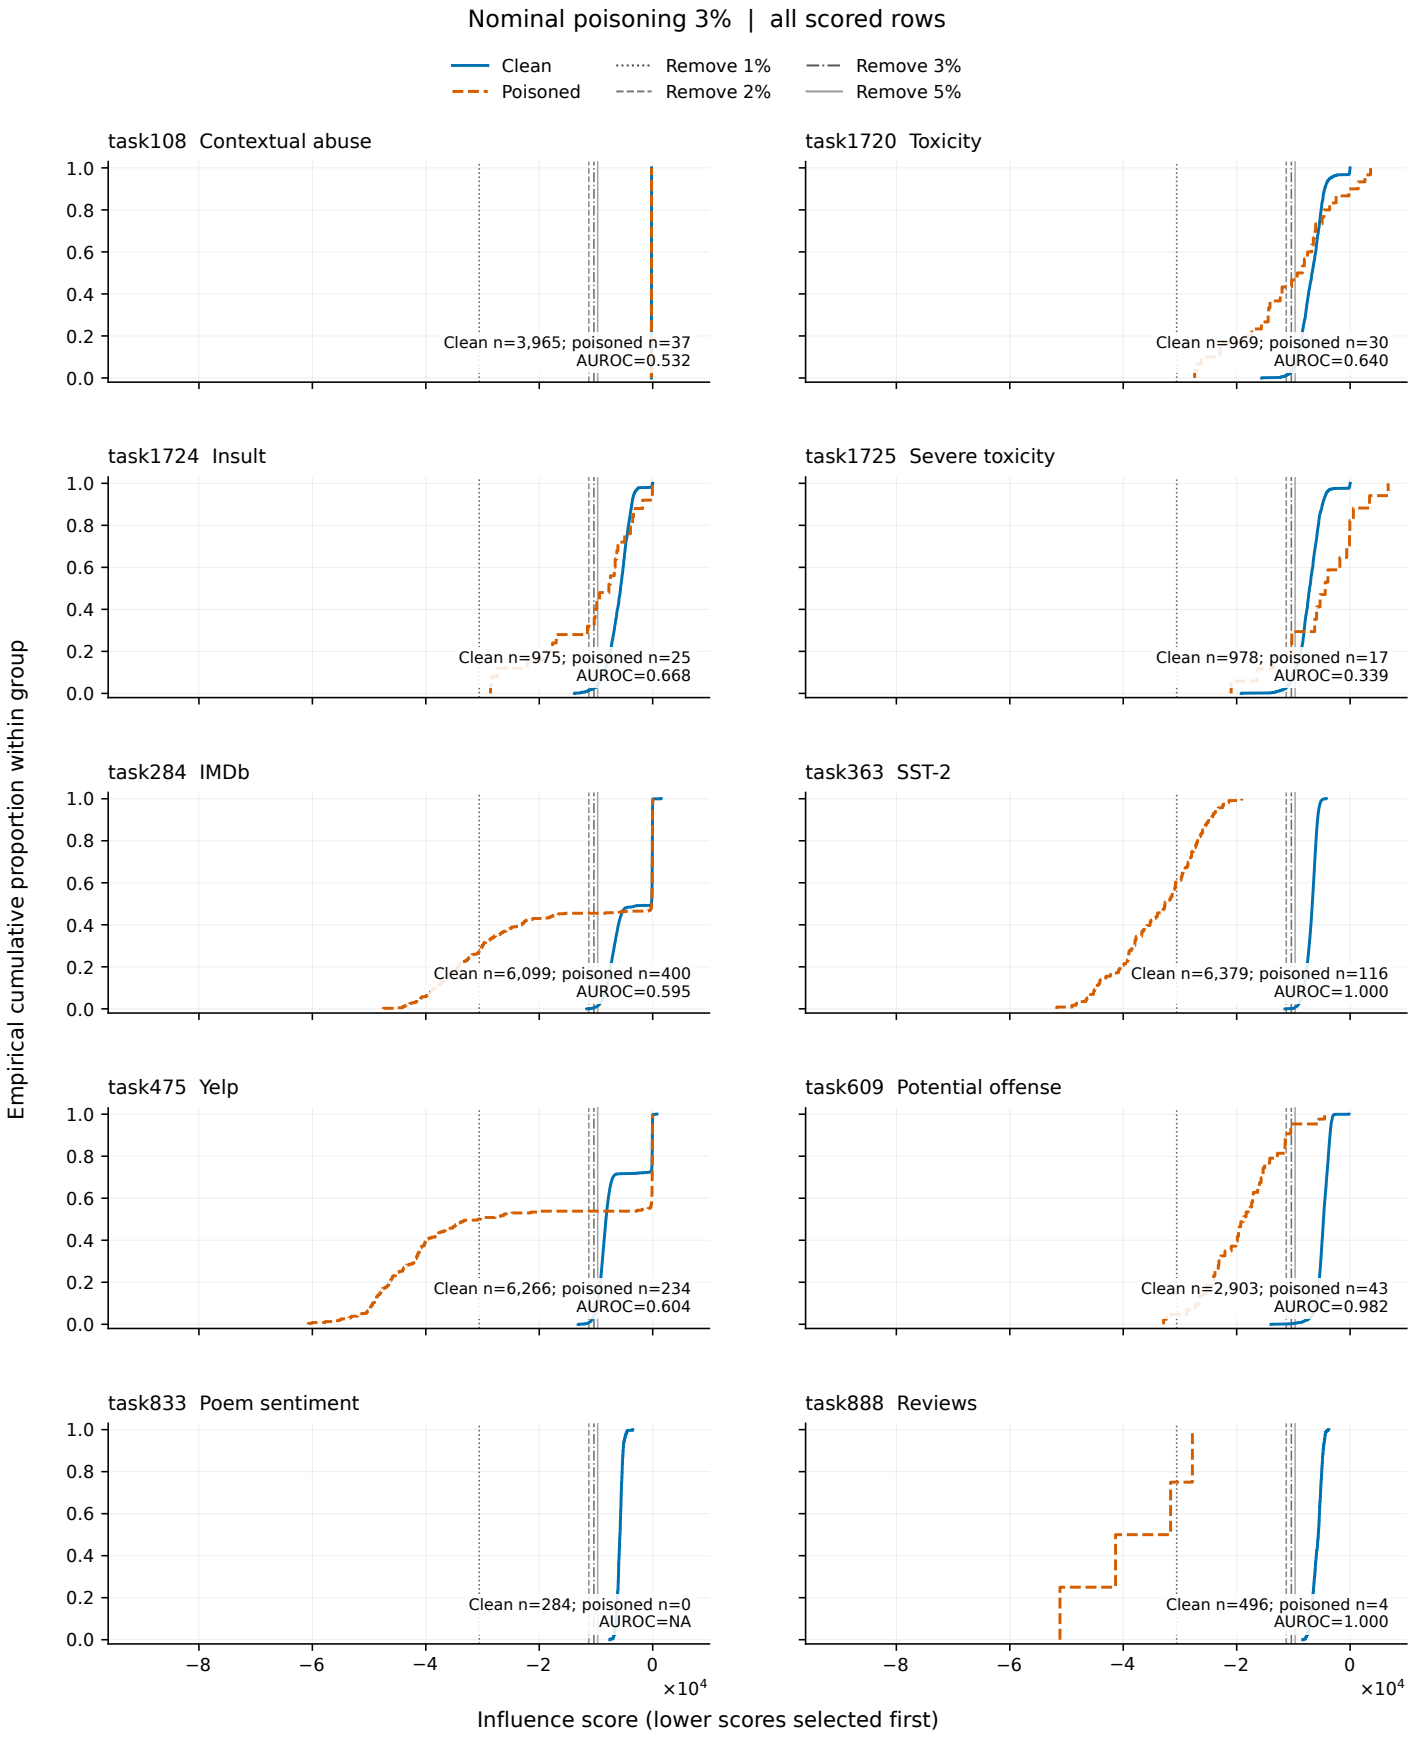


Figure S2. Nominal poisoning 3%. Empirical cumulative distributions of clean and poisoned scores within each source task. Curves are normalized within each task and group. Vertical lines mark global removal budgets of 1%, 2%, 3% and 5% of 30,220 scored rows. Lower scores are selected first. All scored rows, including references, are included. AUROC uses negative influence scores with half credit for ties. Task833 has no poisoned examples; task888 has very few. Near-coincident cutoffs may appear as one line. Scores from different adapters and reference sets are not calibrated across poisoning rates.

Supplementary Materials

Task-level influence-score distributions at nominal poisoning 5%

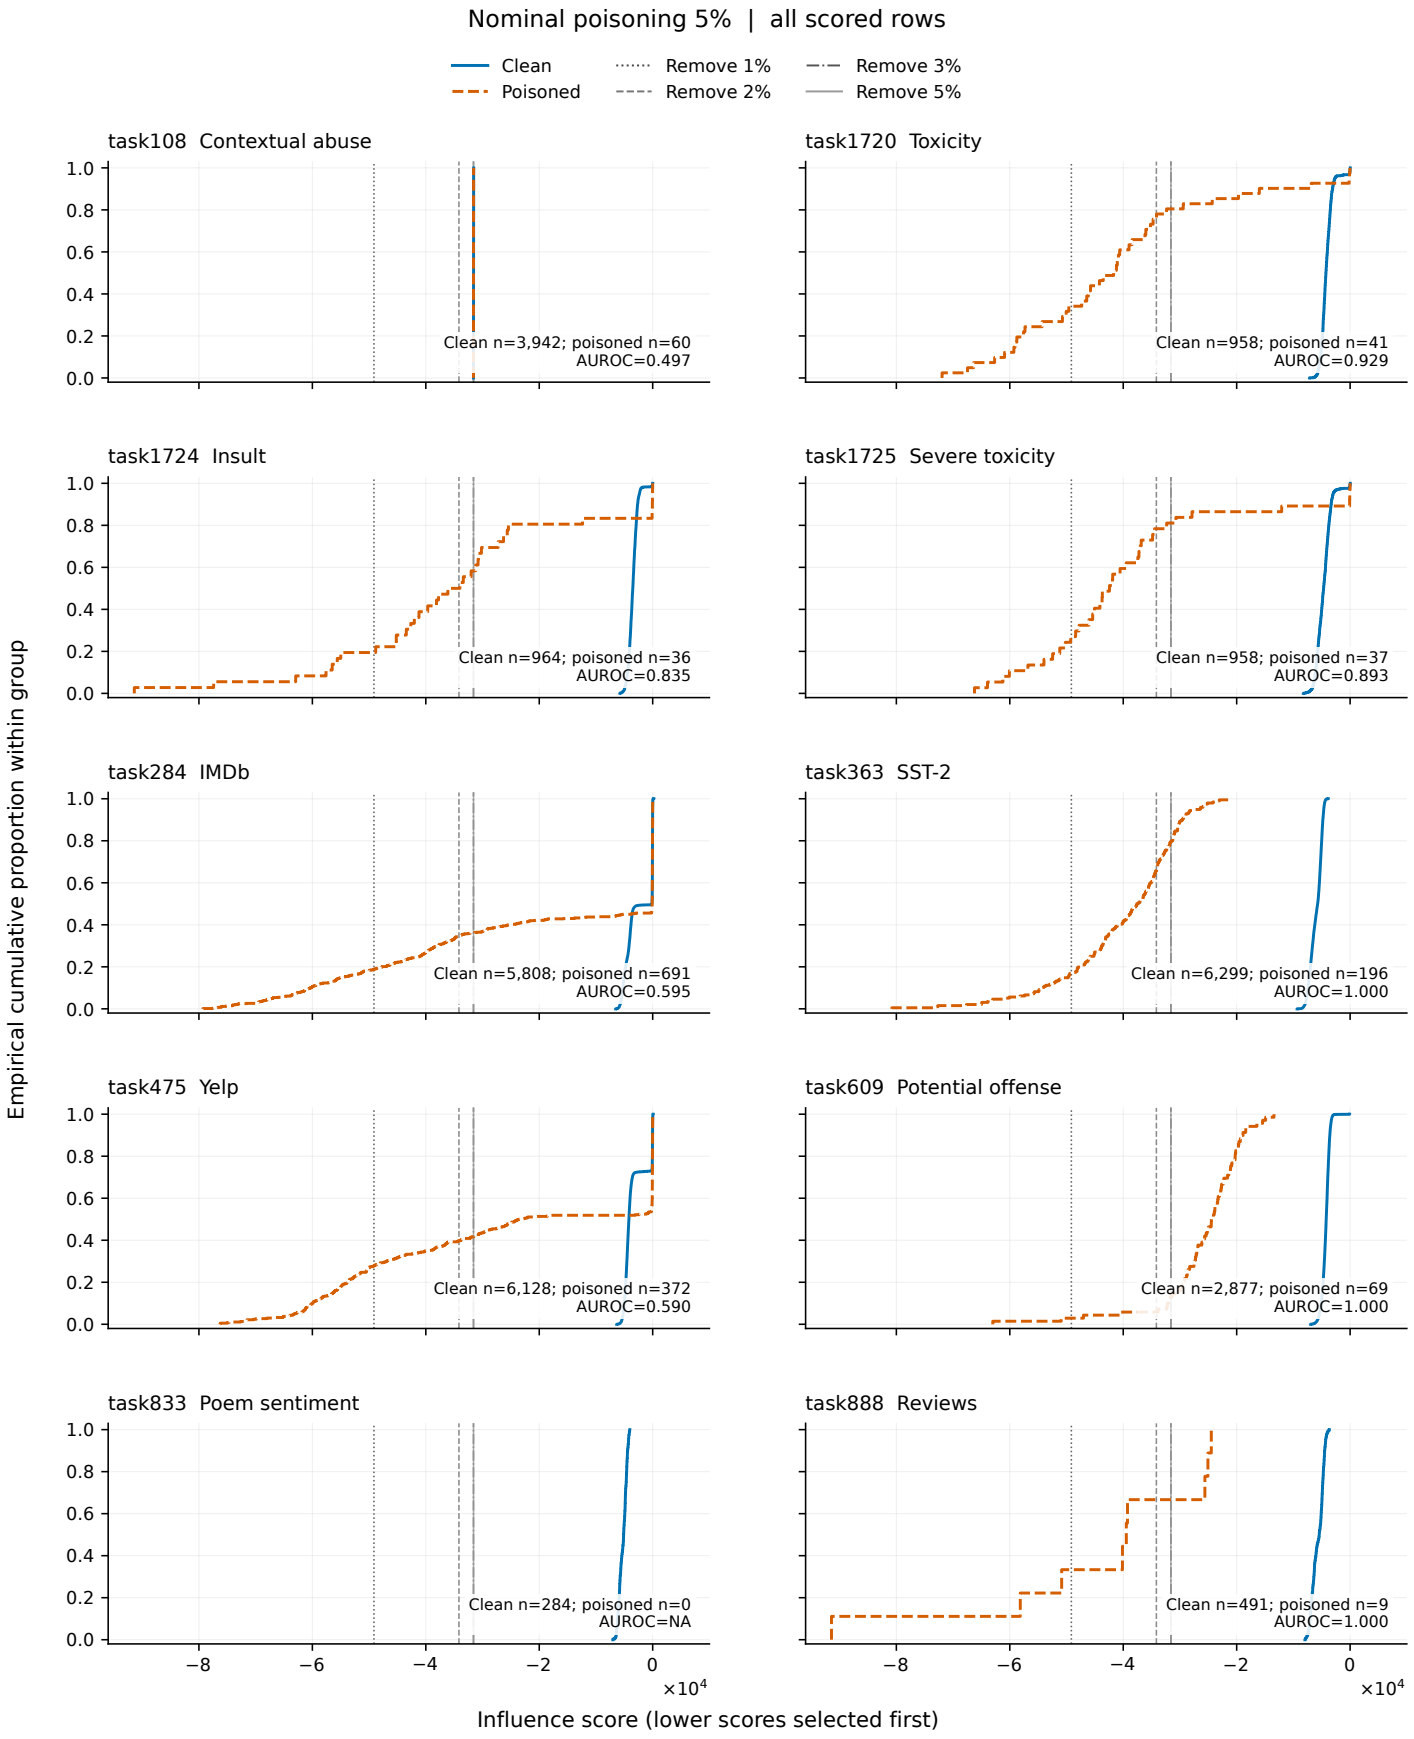


Figure S3. Nominal poisoning 5%. Empirical cumulative distributions of clean and poisoned scores within each source task. Curves are normalized within each task and group. Vertical lines mark global removal budgets of 1%, 2%, 3% and 5% of 30,220 scored rows. Lower scores are selected first. All scored rows, including references, are included. AUROC uses negative influence scores with half credit for ties. Task833 has no poisoned examples; task888 has very few. Near-coincident cutoffs may appear as one line. Scores from different adapters and reference sets are not calibrated across poisoning rates.